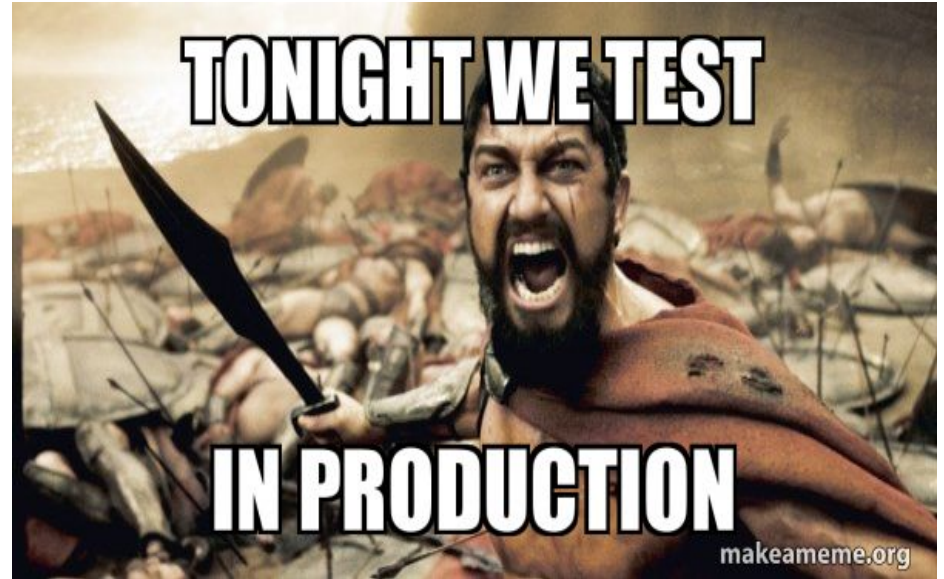


Testing in Production

Diego Pacheco



About me...



- ❑ Cat's Father
- ❑ Head of Software Architect
- ❑ Agile Coach
- ❑ SOA/Microservices Expert
- ❑ DevOps Practitioner
- ❑ Speaker
- ❑ Author



diegopacheco



@diego_pacheco

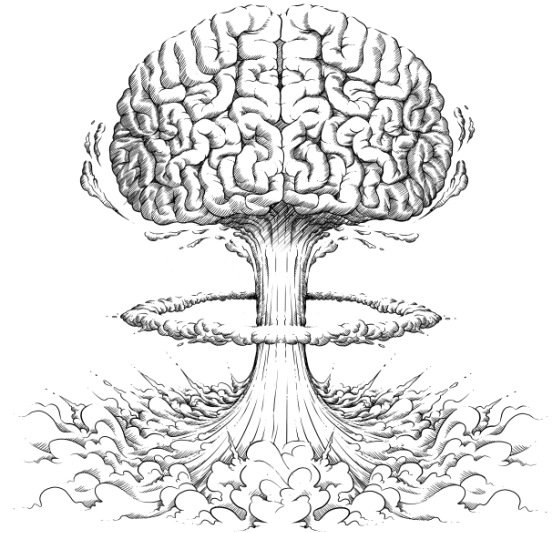


<http://diego-pacheco.blogspot.com.br/>



<https://diegopacheco.github.io/>

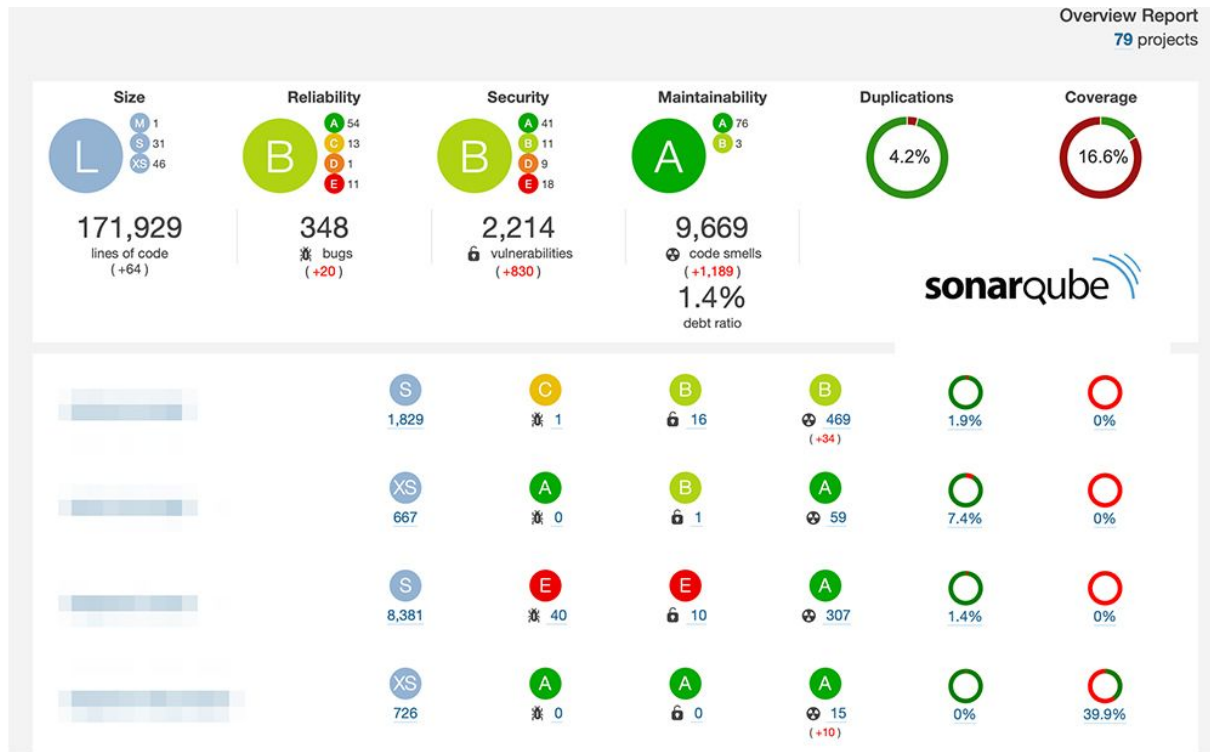
Did you lost your mind?



The issues with pre-production testing - Prediction



100% test coverage is a flawed idea



100% test coverage is a flawed idea

- ❑ Coverage tools just capture (unit and sometimes integration tests) what about the rest?
- ❑ Unit Tests have cost(Mostly Mocks)
- ❑ There is waste in tests. Numbers != Quality.
- ❑ There is no clear and generic/global number about distribution of tests (%)
- ❑ People want signals but there is no such metric can work(plus you will fight configs).



Programming Wisdom @CodeWisdom · Feb 5



"Testing shows the presence, not the absence of bugs." – Edsger W Dijkstra



16



497



2.3K



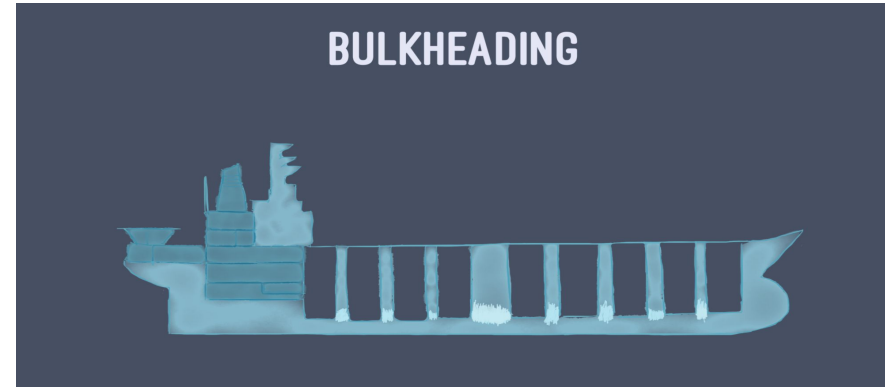
The issues with pre-production testing - Replication



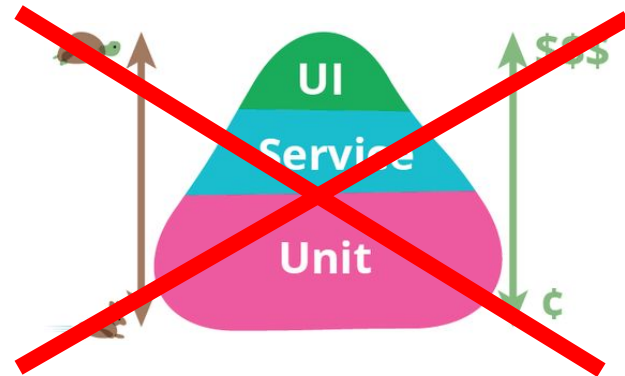
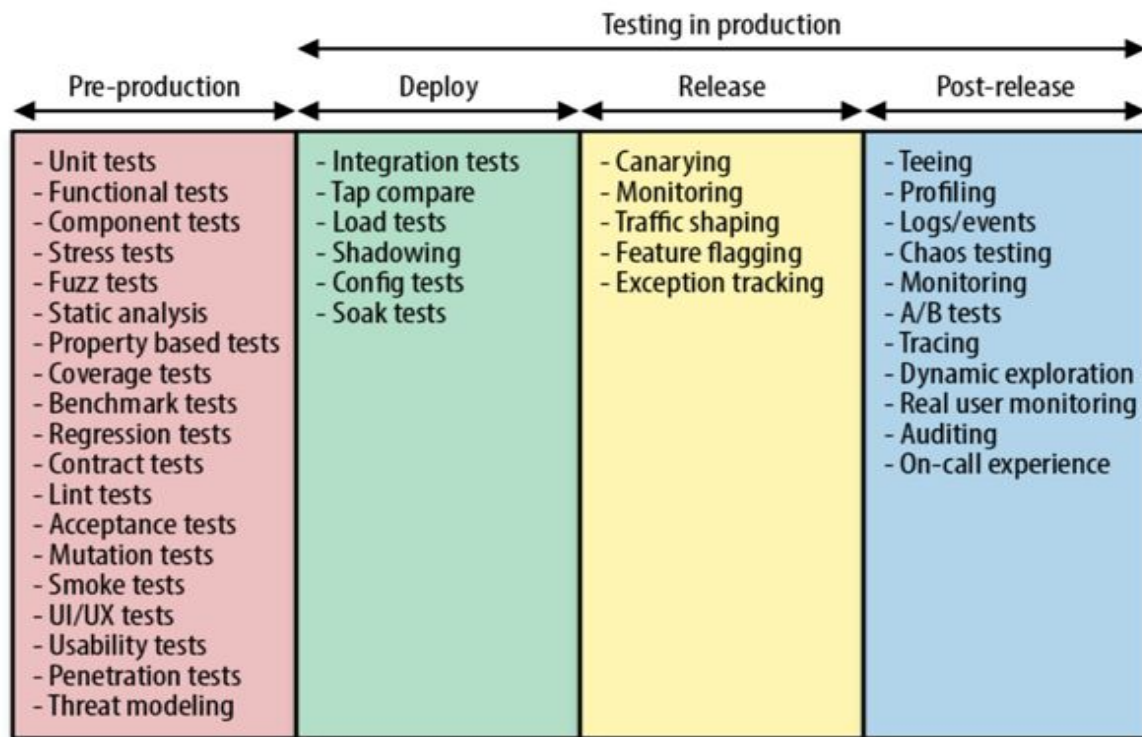
Failure Cannot be legislated/regulated!



- ❑ We need a new approach
- ❑ We need a new mentality
- ❑ Instead of trying to predict everything(which is impossible)
- ❑ We need to have better understand on what's going on and be able to isolate failure and test in production.

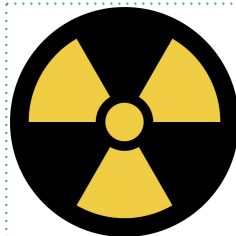


Testing Diversity

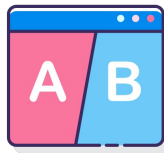
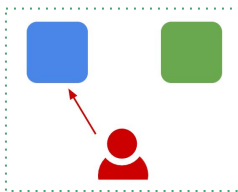


Pattern Landscape

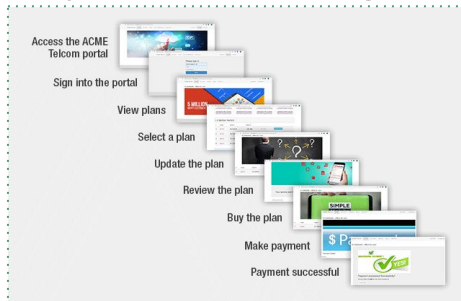
Nuclear Options



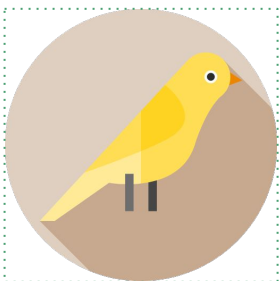
Blue/Green Deployment



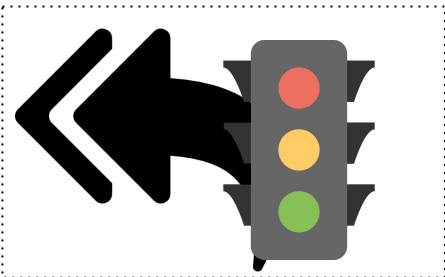
Synthetic monitoring



Automated Canary /
Split Traffic



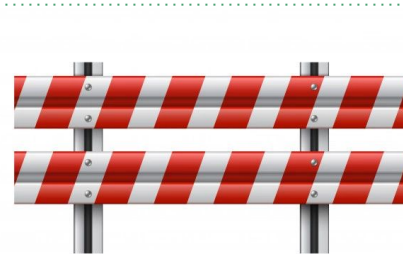
Reply Traffic



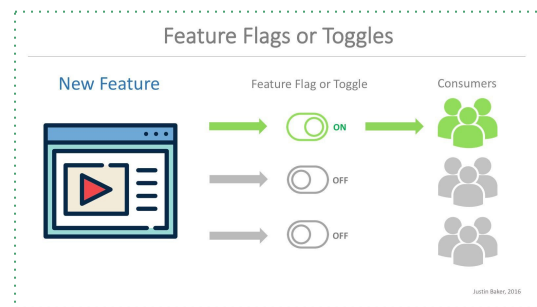
Dark Canary / Dark Launch



Safe Guards

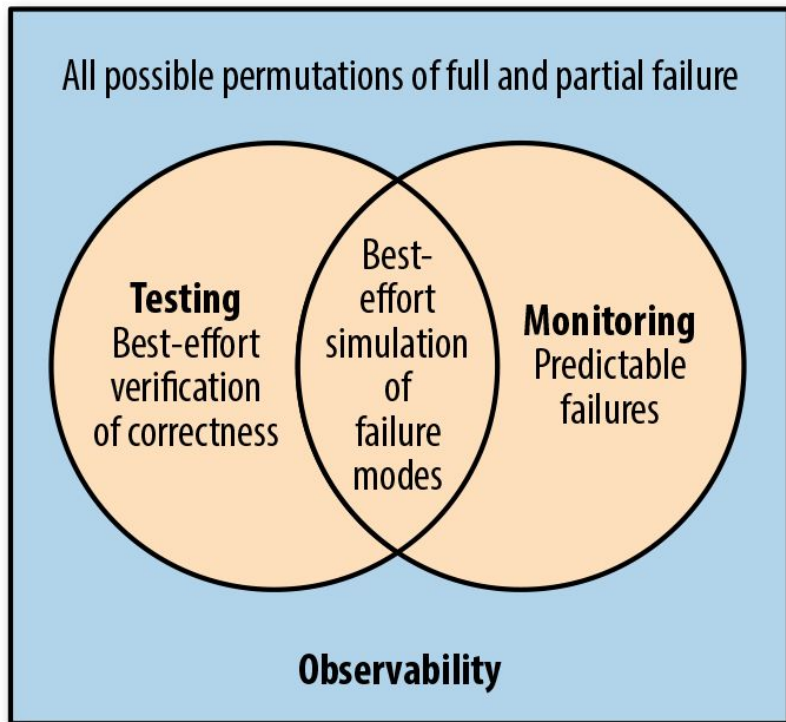
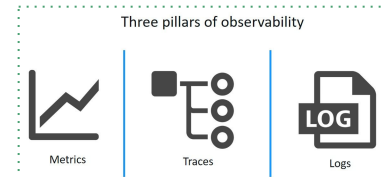


Pattern Landscape - Toggles



- ❑ Common in dramatic new versions (e.g. New News Feed)
- ❑ In a nutshell is Properties + IFs in the Code
- ❑ Can be Role based, Strategy Based. E.g all types of user, from a country, with same age or with x,y,z prop.
- ❑ Require Code Changes (Better than branching)
- ❑ Use the prefix "TEMP" in front of the flag to express it will be clean up soon. E.g "TEMP-feature-flag-01-login"
- ❑ And if you use branches name the branch. E.g "cleanup/TEMP-feature-flag-01-login"
- ❑ Some people love it other hate(abuses and tech debt)
- ❑ Flags should be used with proper Observability.
- ❑ Flag Status - e.g is the flag is not used for more than 7 days is safe to remove.
- ❑ Tech vs Non-Tech Flags (Business Feature vs Migration)

Pattern Landscape - Observability



- ❑ Basic Requirement for everything
- ❑ Whats Going On right now?
- ❑ Is the services working properly?
- ❑ Running your tests against prod is a form of Observability
- ❑ Maturity Levels
 - ❑ 5 - Self-Healing / Self-Operating Systems
 - ❑ 4- Alerts
 - ❑ 3- Dashboards
 - ❑ 2 - Logs
 - ❑ 1 - Praying that all is good
- ❑ Levels
 - ❑ Business - Domain Observability e.g Sales, Logins
 - ❑ Infrastructure e.g k8s pods, Network Traffic
 - ❑ Application e.g RPS, Errors, Tail Latency, queue len.
 - ❑ OS/Machine e.g CPU, Memory, Disk, Network

Pattern Landscape - Synthetic Monitoring

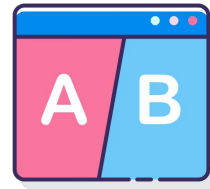


- ❑ A.K.A Semantic Monitoring or simply Business Monitoring
- ❑ Replicates the user experience (Great to measure user Latency)
- ❑ Answer basic and yet important question - is the business working fine? Can the users use the product?
- ❑ Similar to E2E Test or UI Testing but Running in prod and generating metrics/dashboards/alerts.
- ❑ Might Require Instrumentation in other services (e.g make some records invisible to services).
- ❑ Can be done with same E2E/UI Tools such as Selenium or Cypress.
- ❑ Ideally should run outside of your cloud (if you want know the user latency).
- ❑ Like any good test need to be self-contained and have proper isolation.



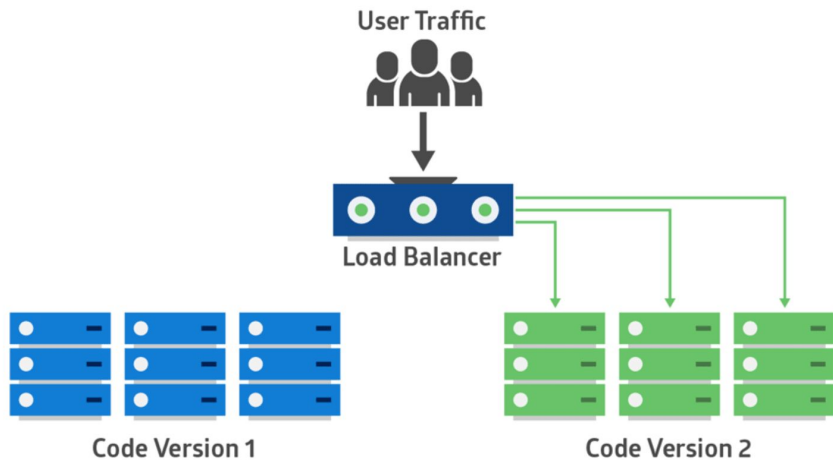
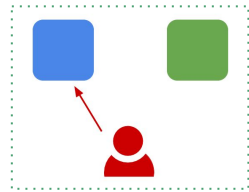
Selenium  cypress.io

A/B/n Testing

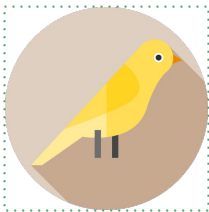


- ❑ Split Testing
- ❑ Testing in Production for Business
- ❑ Used by all Big Tech companies in SV
- ❑ Huge source of Revenue tuning
- ❑ From simple colors and wording to business behaviors.
- ❑ Can be done with simple duplication on the UI. e.g page1 and page2.
- ❑ Easily will require feature toggles and coordination with backend services
- ❑ Microsoft Bing has 30B possible experiments (is very likely 2 users never get the same site).

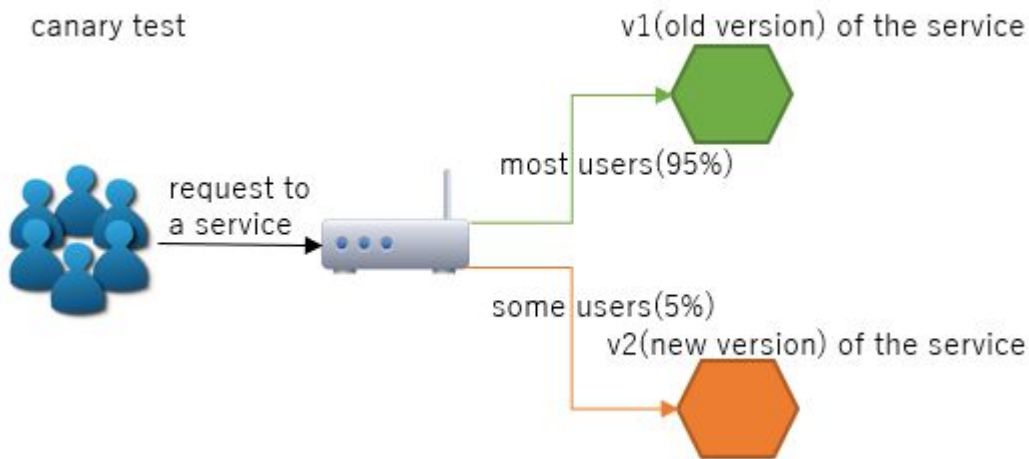
Blue / Green Deployment



- ❑ It's like a switch
- ❑ It's about fast recovery
- ❑ Also reduce the blast radius
- ❑ Unfortunately is all or nothing
- ❑ Also can be used to test things in prod before affecting real users.



Pattern Landscape - Split Traffic / Canary



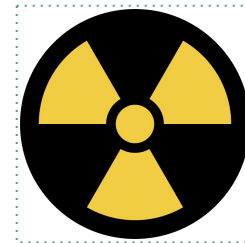
- ❑ IMHO Best form of Testing
- ❑ As real as you can get.
- ❑ It's all about reducing blast radius
- ❑ Works with progressive split traffic 1%, 5%, 10%, 25%, 50%, 75%, 100% of users
- ❑ Can be used with Green/Blue as well
- ❑ Require tooling:
 - ❑ Asgard/Spinnaker
 - ❑ Harness
 - ❑ Istio/Flagger
 - ❑ Argo

Pattern Landscape - Chaos Engineering

- ☐ You cannot legislate failure.
- ☐ Form of testing that make sure your infrastructure in antifragile
- ☐ Based on Experiments and Expectations
- ☐ Netflix runs in prod for years without issues
- ☐ Anti-Fragility mindset - Make sure you can recover rather than “prevent failure”
- ☐ Induce failure:
 - ☐ Tear down one machine
 - ☐ Burn CPU, Disk, IO, Network, Memory
 - ☐ Never Return a service call
 - ☐ Return all nulls
 - ☐ Return String bombs
 - ☐ Do GC Bombs
 - ☐ Introduce Lagging, Lose packages
 - ☐ Tear down a AZ, Tear down a Region



Facebook “Nuclear Options”



- ❑ Testing in Production and Chaos Engineering meet UX
- ❑ Facebook - Change Users Story - e.g Disable Recommendations, chat, Simple feed.
- ❑ UX Degradation -- Netflix
- ❑ Resilience Matrix

Resilience Matrix

Resiliency Matrix			
	Checkout	Admin	Storefront
MySQL Shard	Unavailable	Unavailable	Degraded
MySQL Master	Available	Unavailable	Available
Kafka	Available	Degraded	Available
External HTTP API	Degraded	Available	Unavailable
redis-sessions	Unavailable	Unavailable	Degraded

Graceful Degradation As a Feature

RELIABILITY THROUGH CHAOS ENGINEERING



Design for Failure

Identify the most critical end user functionality.



Inject Failure

Impact your system to be sure your user experience isn't impacted.



Degrade Gracefully

Plan for non critical functionality not to get in the way.



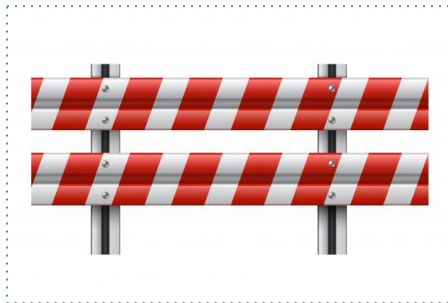
Delight Your Users

Your product metrics will show behaviour, no matter the condition.

Pattern Landscape - Safeguards

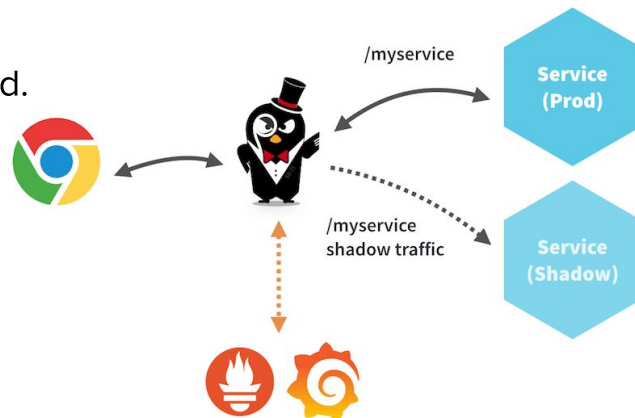
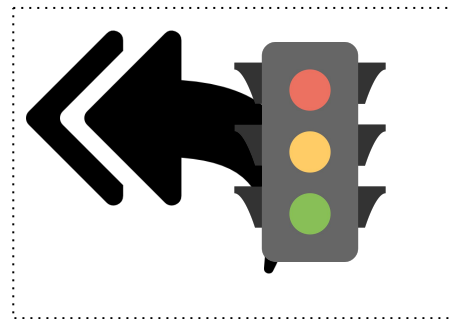
- ❑ Have threshold to STOP.
- ❑ Synthetical monitoring and alerts
- ❑ Make sure Tests are self-contained (Running tests in prod)
- ❑ Observability is the key for everything
- ❑ Dark Pool either in NO-Persist mode or persist in different DB or other env. (Code Instrumentation)
- ❑ Beta Users Program is a great form of isolation and will work well with split traffic (is you have a lot)
- ❑ Traffic split to 1 user first (your internal users in prod maybe?)
- ❑ Use Feature toggles in order to easily and fast switch things off or back to normal.

Safe Guards

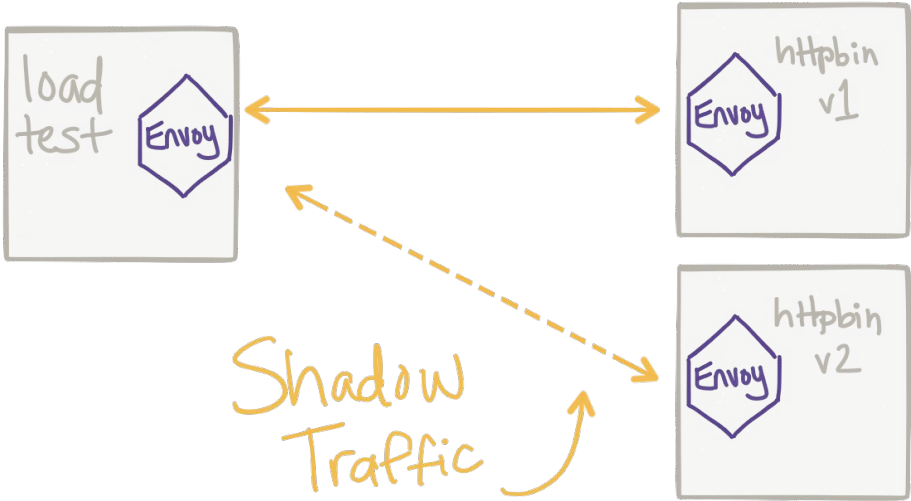
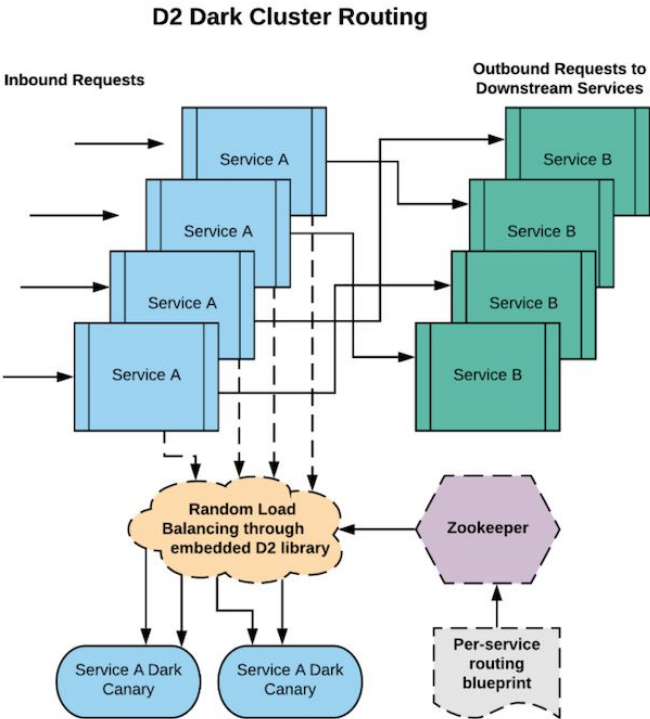


Pattern Landscape - Reply Traffic

- ❑ Used in conjunction with Dark Canary
- ❑ Also known as shadowing or mirroring or dark traffic testing
- ❑ Capture production traffic
- ❑ It's about isolation
- ❑ Should be Zero impact on real users (dark pool)
- ❑ Either by logs or distributed log solution like kafka
- ❑ Reply the traffic in prod(dark canary) or even in non-prod env.
- ❑ Might require some sudo anonymization as move outside of prod.
- ❑ Can get complicated and might require lots of infrastructure



Pattern Landscape - Dark Canary



LinkedIn Dark Canary Formula

Ratio of DarkRequests = $\frac{(\# \text{ DarkServiceACluster instances})}{(\# \text{ ServiceACluster instances})} * \text{multiplier}$

Outbound DarkRequest QPS = (Ratio of DarkRequests) * inbound QPS

$2 \text{ QPS} = \frac{2 \text{ instances}}{100 \text{ instances}} * 1 * 100 \text{ QPS}$

<https://engineering.linkedin.com/blog/2020/production-testing-with-dark-canaries>

Testing in Production

Diego Pacheco

